

Question – 1:

Load the diabetes dataset (assuming it is available in the current working directory)

```
diabetes <- read.csv("diabetes.csv")
```

(a) Number of rows and columns

```
num_rows <- nrow(diabetes)
```

```
num_cols <- ncol(diabetes)
```

```
cat("Number of rows:", num_rows, "\n")
```

```
cat("Number of columns:", num_cols, "\n")
```

(b) Column names

```
col_names <- names(diabetes)
```

```
cat("Column names:", col_names, "\n")
```

(c) Average glyhb

```
average_glyhb <- mean(diabetes$glyhb)
```

```
cat("Average glyhb:", average_glyhb, "\n")
```

(d) Column types

```
column_types <- sapply(diabetes, class)
```

```
cat("Column types:", column_types, "\n")
```

(e) Summary of the dataframe

```
summary(diabetes)
```

(f) Access the age column

```
age_column <- diabetes$age
```

(g) Number of people whose age is greater than 40

```
num_people_age_gt_40 <- sum(diabetes$age > 40)
```

```
cat("Number of people with age > 40:", num_people_age_gt_40, "\n")
```

```
# (h) Female diabetic patients of age >= 30
```

```
female_diabetic_age_30 <- subset(diabetes, diabetes$gender == "Female" & diabetes$age >= 30)
```

```
cat("Female diabetic patients with age >= 30:\n")
```

```
print(female_diabetic_age_30)
```

```
# (i) Patients not from Louisa
```

```
patients_not_from_louisa <- subset(diabetes, diabetes$location != "Louisa")
```

```
cat("Patients not from Louisa:\n")
```

```
print(patients_not_from_louisa)
```

```
> # Load the diabetes dataset (assuming it is available in the current working directory)
> diabetes <- read.csv("diabetes.csv")
> # (a) Number of rows and columns
> num_rows <- nrow(diabetes)
> num_cols <- ncol(diabetes)
> cat("Number of rows:", num_rows, "\n")
Number of rows: 403
> cat("Number of columns:", num_cols, "\n")
Number of columns: 24
> # (b) Column names
> col_names <- names(diabetes)
> cat("Column names:", col_names, "\n")
Column names: X id chol stab.glu hdl ratio glyhb location age gender height weight frame bp.1s bp.1d bp.2s bp.2d waist
hip time.ppn insurance fh smoking dm
> # (c) Average glyhb
> average_glyhb <- mean(diabetes$glyhb)
> cat("Average glyhb:", average_glyhb, "\n")
Average glyhb: NA
> # (d) Column types
> column_types <- sapply(diabetes, class)
> cat("Column types:", column_types, "\n")
Column types: integer integer integer integer integer numeric numeric character integer character integer integer chara
cter integer integer integer integer integer integer integer integer integer integer integer character
> # (e) Summary of the dataframe
> summary(diabetes)
      X      id      chol      stab.glu      hdl      ratio      glyhb
Min.   : 1.0   Min.   :1000   Min.   : 78.0   Min.   : 48.0   Min.   : 12.00   Min.   : 1.500   Min.   : 2.68
1st Qu.:101.5   1st Qu.: 4792   1st Qu.:179.0   1st Qu.: 81.0   1st Qu.: 38.00   1st Qu.: 3.200   1st Qu.: 4.38
Median :202.0   Median :15766   Median :204.0   Median : 89.0   Median : 46.00   Median : 4.200   Median : 4.84
Mean   :202.0   Mean   :15978   Mean   :207.8   Mean   :106.7   Mean   : 50.45   Mean   : 4.522   Mean   : 5.59
3rd Qu.:302.5   3rd Qu.:20336   3rd Qu.:230.0   3rd Qu.:106.0   3rd Qu.: 59.00   3rd Qu.: 5.400   3rd Qu.: 5.60
Max.   :403.0   Max.   :41756   Max.   :443.0   Max.   :385.0   Max.   :120.00   Max.   :19.300   Max.   :16.11
NA's   :1      NA's   :1      NA's   :1      NA's   :13

      location      age      gender      height      weight      frame
Length:403      Min.   :19.00   Length:403      Min.   :52.00   Min.   : 99.0   Length:403
Class :character 1st Qu.:34.00   Class :character 1st Qu.:63.00   1st Qu.:151.0   Class :character
Mode  :character Median :45.00   Mode  :character Median :66.00   Median :172.5   Mode  :character
                Mean   :46.85                Mean   :66.02   Mean   :177.6
                3rd Qu.:60.00                3rd Qu.:69.00   3rd Qu.:200.0
                Max.   :92.00                Max.   :76.00   Max.   :325.0
                NA's   :5                  NA's   :1
```

```

      bp.1s      bp.1d      bp.2s      bp.2d      waist      hip      time.ppn
Min.   : 90.0   Min.   : 48.00   Min.   :110.0   Min.   : 60.00   Min.   :26.0   Min.   :30.00   Min.   : 5.0
1st Qu.:121.2   1st Qu.: 75.00   1st Qu.:138.0   1st Qu.: 84.00   1st Qu.:33.0   1st Qu.:39.00   1st Qu.: 90.0
Median :136.0   Median : 82.00   Median :149.0   Median : 92.00   Median :37.0   Median :42.00   Median : 240.0
Mean   :136.9   Mean   : 83.32   Mean   :152.4   Mean   : 92.52   Mean   :37.9   Mean   :43.04   Mean   : 341.2
3rd Qu.:146.8   3rd Qu.: 90.00   3rd Qu.:161.0   3rd Qu.:100.00   3rd Qu.:41.0   3rd Qu.:46.00   3rd Qu.: 517.5
Max.   :250.0   Max.   :124.00   Max.   :238.0   Max.   :124.00   Max.   :56.0   Max.   :64.00   Max.   :1560.0
NA's    :5      NA's    :5      NA's    :262    NA's    :2      NA's    :2      NA's    :3
insurance      fh      smoking      dm
Min.   :0.000   Min.   :0.0000   Min.   :1.000   Length:403
1st Qu.:0.000   1st Qu.:0.0000   1st Qu.:1.000   Class :character
Median :1.000   Median :0.0000   Median :2.000   Mode  :character
Mean   :1.055   Mean   :0.1886   Mean   :1.846
3rd Qu.:2.000   3rd Qu.:0.0000   3rd Qu.:2.000
Max.   :2.000   Max.   :1.0000   Max.   :3.000

> # (f) Access the age column
> age_column <- diabetes$age
> # (g) Number of people whose age is greater than 40
> num_people_age_gt_40 <- sum(diabetes$age > 40)
> cat("Number of people with age > 40:", num_people_age_gt_40, "\n")
Number of people with age > 40: 243
> # (h) Female diabetic patients of age >= 30
> female_diabetic_age_30 <- subset(diabetes, diabetes$gender == "Female" & diabetes$age >= 30)
> cat("Female diabetic patients with age >= 30:\n")
Female diabetic patients with age >= 30:
> print(female_diabetic_age_30)
[1] x      id      chol      stab.glu      hdl      ratio      glyhb      location      age      gender      height
[12] weight      frame      bp.1s      bp.1d      bp.2s      bp.2d      waist      hip      time.ppn      insurance      fh
[23] smoking      dm
<0 rows> (or 0-length row.names)
> # (i) Patients not from Louisa
> patients_not_from_louisa <- subset(diabetes, diabetes$location != "Louisa")
> cat("Patients not from Louisa:\n")
Patients not from Louisa:
> print(patients_not_from_louisa)
      X      id      chol      stab.glu      hdl      ratio      glyhb      location      age      gender      height      weight      frame      bp.1s      bp.1d      bp.2s      bp.2d      waist
1      1 1000 203      82 56      3.6      4.31 Buckingham 46 female      62      121 medium 118      59      NA      NA      29
2      2 1001 165      97 24      6.9      4.44 Buckingham 29 female      64      218 large 112      68      NA      NA      46
3      3 1002 228      92 37      6.2      4.64 Buckingham 58 female      61      256 large 190      92      185      92      49
4      4 1003 78      93 12      6.5      4.63 Buckingham 67 male      67      119 large 110      50      NA      NA      33

      hip      time.ppn      insurance      fh      smoking      dm
1      38      720      1 0      3 no
2      48      360      0 0      2 no
3      57      180      2 0      2 no
4      38      480      1 0      3 no
5      41      300      0 0      3 yes
6      42      195      1 0      1 no
7      49      720      2 1      2 no
8      39      1020      0 0      2 no
9      40      300      2 1      1 no
10     50      240      2 0      2 no
20     51      540      0 0      2 yes
21     41      1020      2 0      2 no
22     46      480      0 0      2 no
23     55      240      1 1      2 yes
24     44      285      1 0      1 no
25     38      210      2 0      1 no
26     36      780      1 0      2 no
27     47      420      1 1      2 no
28     33      510      0 1      1 no
29     38      720      2 0      3 no
31     40      540      1 0      2 no
32     35      780      1 0      1 no
33     41      225      0 0      2 yes
34     43      330      0 0      3 no
35     34      720      1 0      2 no
36     44      780      2 1      2 no
37     35      720      1 0      1 no
38     49      240      2 0      2 no
39     36      30      0 0      3 no
40     48      840      2 0      3 yes
41     40      720      1 0      2 no

```

output isn't shown completely here... execute it!

Question – 2:

Create a 2x3 matrix M with randomly selected values between 10 and 30

```
M <- matrix(sample(10:30, 6, replace = TRUE), nrow = 2, byrow = TRUE)
```

Create a vector V with three values

```
V <- c(20, 25, 30)
```

Add V as the third row of M

```
M <- rbind(M, V)
```

Print the matrix M

```
print(M)
```

```
> source("~/active-rstudio-document")  
  [,1] [,2] [,3]  
    13  12  26  
    27  29  18  
V    20  25  30
```

Question – 3:

Create a numeric vector

```
x <- c(1, 2, 3, 4, 5)
```

Add zeros at starting and ending of x

```
x <- c(0, 0, x, 0, 0)
```

Print the vector x

```
print(x)
```

```
> source("~/active-rstudio-document")  
[1] 0 0 1 2 3 4 5 0 0
```

Question – 4:

Load the iris dataset

```
library(datasets)
```

```
data(iris)
```

Set the number of clusters

```
k <- 3
```

Initialize the cluster centroids

```
centroids <- kmeans(iris[, 1:4], centers = k, nstart = 10)$centers
```

Assign each data point to a cluster

```
cluster_assignments <- kmeans(iris[, 1:4], centers = centroids)$cluster
```

#print(cluster_assignments)

Plot the data points with their cluster assignments

```
plot(iris[, 1:4], col = c("red", "green", "blue")[cluster_assignments], pch = 19)
```

```
points(centroids, pch = 1, cex = 1.5, col = "black")
```

